

```
# import necessary libraries
import pandas as pd
import numpy as np
from numpy import array
from numpy import asarray
from numpy import zeros

import nltk
nltk.download('stopwords')
nltk.download('averaged_perceptron_tagger')
nltk.download('wordnet')
from nltk.corpus import stopwords
from nltk.stem.wordnet import WordNetLemmatizer
en_stop = set(nltk.corpus.stopwords.words('english'))

from keras.preprocessing.text import one_hot
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers.core import Activation, Dropout, Dense
from keras.layers import Flatten, LSTM
from keras.layers import GlobalMaxPooling1D
from keras.models import Model
from keras.layers.embeddings import Embedding
from sklearn.model_selection import train_test_split
from keras.preprocessing.text import Tokenizer
from keras.layers import Input
from keras.layers.merge import Concatenate

import re
import pickle
import matplotlib.pyplot as plt

from keras import backend as K
from keras.models import load_model
```

 [nltk_data] Downloading package stopwords to /root/nltk_data...
 [nltk_data] Package stopwords is already up-to-date!
 [nltk_data] Downloading package averaged_perceptron_tagger to
 [nltk_data] /root/nltk_data...
 [nltk_data] Unzipping taggers/averaged_perceptron_tagger.zip.
 [nltk_data] Downloading package wordnet to /root/nltk_data...
 [nltk_data] Unzipping corpora/wordnet.zip.
 Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tensorflow1 magic.

```
# read the data in from the csv
movies = pd.read_csv("movies_small_subset_df.csv")
movies = movies[['MovieID', 'MovieName', 'Genre', 'Plot',
                 'clean_plot_text']]
```

```
# Function for converting genre column to a list such that it can be indexed to get g
def format_list(x):
    x = x.replace("'", "")
    x = x.replace("[", "")
    x = x.replace("]", "")
    x = x.split(',')
    result = []
    for word in x:
        result.append(word.strip())
    return result
movies["Genre"] = movies["Genre"].apply(lambda x : format_list(x))
movies.head(5)
```

	MovieID	MovieName	Genre	Plot	clean_plot_text
0	23890098	Taxi Blues	[Drama, World cinema]	Shlykov, a hard-working taxi driver and Lyosha...	shlykov hard work taxi driver lyosha saxophoni...
1	31186339	The Hunger Games	[Action, Drama]	The nation of Panem consists of a wealthy Capi...	nation panem consist wealthy capitol twelve po...
2	20663735	Narasimham	[Action,	Poovalli Induchoodan is	poovalli induchoodan sentence six year

```
movies.shape
```

```
(35523, 5)
```


```
# get a list of all unique genres from genres column in dataframe
unique_genre_list = list(set([a for b in movies.Genre.tolist() for a in b]))

for i in range(len(unique_genre_list)):
    movies[unique_genre_list[i]] = pd.Series([0 for x in range(len(movies.index))], ind


movies.head(2)
```

	MovieID	MovieName	Genre	Plot	clean_plot_text	Short Film	Comedy	Romance Film	T
0	23890098	Taxi Blues	[Drama, World cinema]	Shlykov, a hard-working taxi driver and Lyosha...	shlykov hard work taxi driver lyosha saxophoni...	0	0	0	

```
movies.shape
```

 (35523, 13)

```
for gen in unique_genre_list:
    movies[gen] = movies["Genre"].apply(lambda x : (pd.Series([gen]).isin(x)).astype(int))
movies.head(5)
```




	MovieID	MovieName	Genre	Plot	clean_plot_text	Short Film	Comedy	Romance
0	23890098	Taxi Blues	[Drama, World cinema]	Shlykov, a hard-working taxi driver and Lyosha...	shlykov hard work taxi driver lyosha saxophoni...	0	0	
1	31186339	The Hunger Games	[Action, Drama]	The nation of Panem consists of a wealthy Capi...	nation panem consist wealthy capitol twelve po...	0	0	
2	20663735	Narasimham	[Action, Drama]	Poovalli Induchoodan is sentenced for six yea...	poovalli induchoodan sentence six year prison ...	0	0	
				The Lemon				

```
#movies.to_csv("movies_one_hot_df.csv")
#movies = pd.read_csv("movies_one_hot_df.csv")

sentences = list(movies["clean_plot_text"])
sentence_list = [ sen.split(' ') for sen in sentences]
plt.hist([len(s) for s in sentence_list], bins=50)
plt.xlabel('Count of words')
plt.ylabel('Frequency of Sentences')

plt.show()
```





```

X = []
for sen in sentences:
    X.append(sen)
# Create output set (target/labels)
y = movies[unique_genre_list].values

# Split it into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_stat

X_train1 = list(str(elem) for elem in X_train)
X_test1 = list(str(elem) for elem in X_test)

tokenizer_movie = Tokenizer(num_words=5000)
tokenizer_movie.fit_on_texts(X_train1)

# saving
with open('tokenizer_movie_final.pickle', 'wb') as handle:
    pickle.dump(tokenizer_movie, handle, protocol=pickle.HIGHEST_PROTOCOL)
# loading
with open('tokenizer_movie_final.pickle', 'rb') as handle:
    tokenizer_movie = pickle.load(handle)

X_train1 = tokenizer_movie.texts_to_sequences(X_train1)
X_test1 = tokenizer_movie.texts_to_sequences(X_test1)

vocab_size = len(tokenizer_movie.word_index) + 1

maxlen = 500

X_train1 = pad_sequences(X_train1, padding='post', maxlen=maxlen)
X_test1 = pad_sequences(X_test1, padding='post', maxlen=maxlen)

```

```

# Define helper functions to get pre-trained glove word vector embeddings
# and create an embeddings matrix

def get_word_embeddings():
    embeddings_dictionary = dict()
    glove_file = open('glove.6B.100d.txt', encoding="utf8")

    for line in glove_file:
        records = line.split()
        word = records[0]
        vector_dimensions = asarray(records[1:], dtype='float32')
        embeddings_dictionary[word] = vector_dimensions

    glove_file.close()
    return embeddings_dictionary

embeddings_dictionary = get_word_embeddings()

def get_embedding_matrix():
    embedding_matrix = zeros((vocab_size, 100))
    for word, index in tokenizer_movie.word_index.items():
        embedding_vector = embeddings_dictionary.get(word)
        if embedding_vector is not None:
            embedding_matrix[index] = embedding_vector
    return embedding_matrix

embedding_matrix = get_embedding_matrix()

# Define functions to be able to calculate additional metrics like precision, recall,

def recall_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    possible_positives = K.sum(K.round(K.clip(y_true, 0, 1)))
    recall = true_positives / (possible_positives + K.epsilon())
    return recall

def precision_m(y_true, y_pred):
    true_positives = K.sum(K.round(K.clip(y_true * y_pred, 0, 1)))
    predicted_positives = K.sum(K.round(K.clip(y_pred, 0, 1)))
    precision = true_positives / (predicted_positives + K.epsilon())
    return precision

def f1_m(y_true, y_pred):
    precision = precision_m(y_true, y_pred)
    recall = recall_m(y_true, y_pred)
    return 2*((precision*recall)/(precision+recall+K.epsilon()))

def hamming_loss(y_true, y_pred):
    return K.mean(y_true*(1-y_pred)+(1-y_true)*y_pred)

```

```
# Approach 1
# Use a single dense layer with six outputs with sigmoid activation functions and bin
# Each neuron in the output dense layer will represent one of the six output labels.
# ACTIVATION : SIGMOID
from tensorflow import set_random_seed
set_random_seed(1)
deep_inputs_single = Input(shape=(maxlen,))
embedding_layer_single = Embedding(vocab_size, 100, weights=[embedding_matrix], train
LSTM_Layer_1_single = LSTM(128)(embedding_layer_single)
dense_layer_1_single = Dense(8, activation='sigmoid')(LSTM_Layer_1_single)
model_movie_single = Model(inputs=deep_inputs_single, outputs=dense_layer_1_single)

model_movie_single.compile(loss='binary_crossentropy', optimizer='adam', metrics=['ac

print(model_movie_single.summary())
```

☞ Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 500)	0
=====		
embedding_1 (Embedding)	(None, 500, 100)	10465300
=====		
lstm_1 (LSTM)	(None, 128)	117248
=====		
dense_1 (Dense)	(None, 8)	1032
=====		
Total params: 10,583,580		
Trainable params: 118,280		
Non-trainable params: 10,465,300		
=====		
None		

```
history_movie_single_5 = model_movie_single.fit(X_train1, y_train, batch_size=128, ep
```

☞ Train on 22734 samples, validate on 5684 samples

```
Epoch 1/5
22734/22734 [=====] - 134s 6ms/step - loss: 0.4890 - acc
Epoch 2/5
22734/22734 [=====] - 132s 6ms/step - loss: 0.4746 - acc
Epoch 3/5
22734/22734 [=====] - 132s 6ms/step - loss: 0.4740 - acc
Epoch 4/5
22734/22734 [=====] - 131s 6ms/step - loss: 0.4740 - acc
Epoch 5/5
22734/22734 [=====] - 131s 6ms/step - loss: 0.4750 - acc
```

```
loss, accuracy, f1_score, precision, recall, hamming = model_movie_single.evaluate(X_
```

```

print("Test Score:", loss)
print("Test Accuracy:", accuracy)
print("Test Precision:", precision)
print("Test Recall:", recall)
print("Test F1-score:", f1_score)
print("Test hamming_loss:", hamming)

```

```

☞ 7105/7105 [=====] - 52s 7ms/step
Test Score: 0.4734555327162451
Test Accuracy: 0.7906052076002815
Test Precision: 0.5420741732904057
Test Recall: 0.3094173200989508
Test F1-score: 0.3932593436747852
Test hamming_loss: 0.3098531847642728

```

```

model_movie_single.compile(loss='binary_crossentropy', optimizer='adam', metrics=['ac
history_movie_single_10 = model_movie_single.fit(X_train1, y_train, batch_size=128, e

```

```

☞ Train on 22734 samples, validate on 5684 samples
Epoch 1/10
22734/22734 [=====] - 134s 6ms/step - loss: 0.4735 - acc
Epoch 2/10
22734/22734 [=====] - 132s 6ms/step - loss: 0.4729 - acc
Epoch 3/10
22734/22734 [=====] - 132s 6ms/step - loss: 0.4728 - acc
Epoch 4/10
22734/22734 [=====] - 132s 6ms/step - loss: 0.4730 - acc
Epoch 5/10
22734/22734 [=====] - 132s 6ms/step - loss: 0.4579 - acc
Epoch 6/10
22734/22734 [=====] - 131s 6ms/step - loss: 0.4430 - acc
Epoch 7/10
22734/22734 [=====] - 131s 6ms/step - loss: 0.4395 - acc
Epoch 8/10
22734/22734 [=====] - 132s 6ms/step - loss: 0.4382 - acc
Epoch 9/10
22734/22734 [=====] - 131s 6ms/step - loss: 0.4310 - acc
Epoch 10/10
22734/22734 [=====] - 130s 6ms/step - loss: 0.4203 - acc

```

```

loss, accuracy, f1_score, precision, recall, hamming = model_movie_single.evaluate(X_

```

```

print("Test Score:", loss)
print("Test Accuracy:", accuracy)
print("Test Precision:", precision)
print("Test Recall:", recall)
print("Test F1-score:", f1_score)
print("Test hamming_loss:", hamming)

```

```

☞

```

```

7105/7105 [=====] - 53s 7ms/step
Test Score: 0.4190679441699673
Test Accuracy: 0.811963406052076
Test Precision: 0.6141312826694861
Test Recall: 0.3926639343927143
Test F1-score: 0.47751756736210416
Test hamming_loss: 0.26848946765725834

```

```
model_movie_single.save("movie_lstm_single_10.h5")
```

```
# load model from single file
```

```
movie_lstm_single_10 = load_model('movie_lstm_single_10.h5', custom_objects={'f1_m':
```

```
# Approach 1 with softmax activation
```

```
# Use a single dense layer with six outputs with sigmoid activation functions and bin
```

```
# Each neuron in the output dense layer will represent one of the six output labels.
```

```
# ACTIVATION : SOFTMAX
```

```
from tensorflow import set_random_seed
```

```
set_random_seed(1)
```

```
deep_inputs_single = Input(shape=(maxlen,))
```

```
embedding_layer_single = Embedding(vocab_size, 100, weights=[embedding_matrix], train
```

```
LSTM_Layer_1_single = LSTM(128)(embedding_layer_single)
```

```
dense_layer_1_single = Dense(8, activation='softmax')(LSTM_Layer_1_single)
```

```
model_movie_single_soft = Model(inputs=deep_inputs_single, outputs=dense_layer_1_sing
```

```
model_movie_single_soft.compile(loss='binary_crossentropy', optimizer='adam', metrics
```

```
history_movie_single_5_soft = model_movie_single_soft.fit(X_train1, y_train, batch_si
```

```
☞ Train on 22734 samples, validate on 5684 samples
```

```
Epoch 1/5
```

```
22734/22734 [=====] - 136s 6ms/step - loss: 0.5170 - acc
```

```
Epoch 2/5
```

```
22734/22734 [=====] - 134s 6ms/step - loss: 0.5123 - acc
```

```
Epoch 3/5
```

```
22734/22734 [=====] - 133s 6ms/step - loss: 0.5119 - acc
```

```
Epoch 4/5
```

```
22734/22734 [=====] - 133s 6ms/step - loss: 0.5113 - acc
```

```
Epoch 5/5
```

```
22734/22734 [=====] - 133s 6ms/step - loss: 0.5135 - acc
```

```
loss, accuracy, f1_score, precision, recall, hamming = model_movie_single_soft.evaluat
```

```
print("Test Score:", loss)
```

```
print("Test Accuracy:", accuracy)
```

```
print("Test Precision:", precision)
```

```
print("Test Recall:", recall)
```

```
print("Test F1-score:", f1_score)
```

```
print("Test hamming_loss:", hamming)
```



```
↳ 7105/7105 [=====] - 52s 7ms/step
Test Score: 0.511114438531434
Test Accuracy: 0.7802076002814919
Test Precision: 0.045038699768195266
Test Recall: 0.0007870505973539412
Test F1-score: 0.0015467913112197436
Test hamming_loss: 0.26534659909497343
```