

svm

November 9, 2019

0.1 Support Vector Machine model

```
In [ ]: ##### Import necessary libraries
import pandas as pd
from sklearn import metrics, svm
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.calibration import CalibratedClassifierCV
import pickle
```

```
In [2]: # read the first 500,000 yelp reviews
# df = pd.read_json('yelp_dataset/review.json', lines = True)
# df = df[0:500000]
df = pd.read_csv("yelp_dataset/yelp_reviews.csv", encoding='utf-8')
```

```
In [3]: df.head(5)
```

```
Out [3]:
```

	business_id	cool	date	funny	\
0	ujmEBvifdJM6h6RLv4wQIg	0	2013-05-07 04:34:36	1	
1	NZnhc2sEQy3RmzKTZnqtWQ	0	2017-01-14 21:30:33	0	
2	WTqjgwHlXbSFevF32_DJVw	0	2016-11-09 20:09:03	0	
3	ikCg8xy5JIg_NGPx-MSIDA	0	2018-01-09 20:56:38	0	
4	b1b1eb3uo-w561D0ZfCEiQ	0	2018-01-30 23:07:38	0	

	review_id	stars	\
0	Q1sbwvVQXV2734tPgoKj4Q	1	
1	GJXCdrto3ASJOqKeVWPi6Q	5	
2	2TzJjDVDEuAW6MR5Vuc1ug	5	
3	yiOR0Ugj_xUx_Nek0-_Qig	5	
4	11a8sVPMUFTaC7_ABRkmtw	1	

	text	useful	\
0	Total bill for this horrible service? Over \$8G...	6	
1	I *adore* Travis at the Hard Rock's new Kelly ...	0	
2	I have to say that this office really has it t...	3	
3	Went in for a lunch. Steak sandwich was delici...	0	
4	Today was my second out of three sessions I ha...	7	

```

            user_id
0  hG7b0MtEbXx5QzbzE6C_VA
1  yXQM5uF2jS6es16SJzNHfg
2  n6-Gk65cPZL6Uz8qRm3NYw
3  dacAIZ6fTM6mqwW5uxkskg
4  ssoyf2_x0EQMed6fgHeMyQ

```

```
In [4]: df.describe()
```

```

Out[4]:

```

	cool	funny	stars	useful
count	500000.000000	500000.000000	500000.000000	500000.000000
mean	0.551726	0.453300	3.729382	1.307716
std	2.035998	1.679424	1.455030	2.979647
min	0.000000	0.000000	1.000000	0.000000
25%	0.000000	0.000000	3.000000	0.000000
50%	0.000000	0.000000	4.000000	0.000000
75%	0.000000	0.000000	5.000000	1.000000
max	203.000000	146.000000	5.000000	201.000000

```
In [5]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500000 entries, 0 to 499999
Data columns (total 9 columns):
business_id    500000 non-null object
cool           500000 non-null int64
date           500000 non-null object
funny          500000 non-null int64
review_id      500000 non-null object
stars          500000 non-null int64
text           500000 non-null object
useful         500000 non-null int64
user_id        500000 non-null object
dtypes: int64(4), object(5)
memory usage: 34.3+ MB

```

```
In [3]: # split the dataset into training and validation datasets
```

```
train_x, valid_x, train_y, valid_y = train_test_split(df['text'], df['stars'])
```

```
In [4]: # TRAIN THE MODEL AND CALCULATE PERFORMANCE METRICS (ACCURACY, PRECISION, RECALL, F-Score)
# FOR BOTH TRAINING AND TEST SET
```

```
# Weighted performance metrics
```

```
def train_model(classifier, feature_vector_train, label, feature_vector_valid):
```

```
    # fit the training dataset on the classifier
```

```
    classifier.fit(feature_vector_train, label)
```

```
    # predict the labels on training dataset (to compare performance metrics against training set)
```

```
    train_predictions = classifier.predict(feature_vector_train)
```

```

# predict the labels on test dataset
test_predictions = classifier.predict(feature_vector_valid)

# metrics for training dataset
train_accuracy = metrics.accuracy_score(label, train_predictions)
train_precision = metrics.precision_score(label, train_predictions, average = 'weighted')
train_recall = metrics.recall_score(label, train_predictions, average = 'weighted')
train_f1_score = metrics.f1_score(label, train_predictions, average = 'weighted')

# metrics for test dataset
test_accuracy = metrics.accuracy_score(valid_y, test_predictions)
test_precision = metrics.precision_score(valid_y, test_predictions, average = 'weighted')
test_recall = metrics.recall_score(valid_y, test_predictions, average = 'weighted')
test_f1_score = metrics.f1_score(valid_y, test_predictions, average = 'weighted')

return [test_accuracy, test_precision, test_recall, test_f1_score], [train_accuracy, train_precision, train_recall, train_f1_score]

```

* Note : the `TfidfVectorizer` conducts most of the pre-processing steps such as converting to lower case, removing non alpha numeric characters, removing stop words (using `max_df`). Hence the pre-processing step is not included for logistic regression

```

In [8]: # word level tf-idf
tfidf_vect = TfidfVectorizer(lowercase = True, analyzer='word', token_pattern=r'[a-zA-Z\d]+',
                             max_features=500)

tfidf_vect.fit(df['text'])
xtrain_tfidf = tfidf_vect.transform(train_x)
xvalid_tfidf = tfidf_vect.transform(valid_x)

```

0.2 Model 1 : Bag of word representation - Word level

```

In [9]: # SVM Classifier on Word Level TF IDF Vectors
# C (penalty) = 1
# Gamma = "auto"
# Kernel : RBF (Default)
# Default for max_iter is -1 which means there is no limit to the number of iterations
# First attempt to train the SVM classifier with this default value ran endlessly
# Second attempt is made by setting max_iter to 1000

results = train_model(svm.SVC(gamma = "auto", max_iter = 1000), xtrain_tfidf, train_y,

print ("SVM, WordLevel TF-IDF train accuracy: ", results[1][0])
print("")
print ("SVM, WordLevel TF-IDF train precision: ", results[1][1])
print("")
print ("SVM, WordLevel TF-IDF train recall: ", results[1][2])
print("")
print ("SVM, WordLevel TF-IDF train f1_score: ", results[1][3])
print("*****")

```

```

print ("SVM, WordLevel TF-IDF test accuracy: ", results[0][0])
print("")
print ("SVM, WordLevel TF-IDF test precision: ", results[0][1])
print("")
print ("SVM, WordLevel TF-IDF test recall: ", results[0][2])
print("")
print ("SVM, WordLevel TF-IDF test f1_score: ", results[0][3])

```

```

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning:
  % self.max_iter, ConvergenceWarning)

```

SVM, WordLevel TF-IDF train accuracy: 0.137424

SVM, WordLevel TF-IDF train precision: 0.21459502128943445

SVM, WordLevel TF-IDF train recall: 0.137424

SVM, WordLevel TF-IDF train f1_score: 0.0622646623396362

SVM, WordLevel TF-IDF test accuracy: 0.137904

SVM, WordLevel TF-IDF test precision: 0.2274812152626558

SVM, WordLevel TF-IDF test recall: 0.137904

SVM, WordLevel TF-IDF test f1_score: 0.06328901067423066

```

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:114:
  'precision', 'predicted', average, warn_for)
/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/metrics/classification.py:114:
  'precision', 'predicted', average, warn_for)

```

0.3 Model : Bag of word representation - Word level : Change max_iter

```

In [10]: # SVM Classifier on Word Level TF IDF Vectors
         # C (penalty) = 1
         # Gamma = "auto"
         # Kernel : RBF (Default)
         # Third attempt is made by setting max_iter to 10000
         results = train_model(svm.SVC(gamma = "auto", max_iter = 10000), xtrain_tfidf, train_y)

print ("SVM, WordLevel TF-IDF train accuracy: ", results[1][0])
print("")
print ("SVM, WordLevel TF-IDF train precision: ", results[1][1])
print("")
print ("SVM, WordLevel TF-IDF train recall: ", results[1][2])

```

```

print("")
print ("SVM, WordLevel TF-IDF train f1_score: ", results[1][3])
print("*****")
print ("SVM, WordLevel TF-IDF test accuracy: ", results[0][0])
print("")
print ("SVM, WordLevel TF-IDF test precision: ", results[0][1])
print("")
print ("SVM, WordLevel TF-IDF test recall: ", results[0][2])
print("")
print ("SVM, WordLevel TF-IDF test f1_score: ", results[0][3])

```

```

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning
  % self.max_iter, ConvergenceWarning)

```

SVM, WordLevel TF-IDF train accuracy: 0.18882666666666667

SVM, WordLevel TF-IDF train precision: 0.26546280430132546

SVM, WordLevel TF-IDF train recall: 0.18882666666666667

SVM, WordLevel TF-IDF train f1_score: 0.17675089850410858

SVM, WordLevel TF-IDF test accuracy: 0.189192

SVM, WordLevel TF-IDF test precision: 0.26643551590198383

SVM, WordLevel TF-IDF test recall: 0.189192

SVM, WordLevel TF-IDF test f1_score: 0.1774337201487341

```

In [5]: # ngram level tf-idf
tfidf_vect_ngram = TfidfVectorizer(lowercase = True, analyzer='word', token_pattern=r'
                                     ngram_range=(1,3), max_features=500)

# Fit the model
tfidf_ngram_transformer = tfidf_vect_ngram.fit(df['text'])
xtrain_tfidf_ngram = tfidf_ngram_transformer.transform(train_x)
xvalid_tfidf_ngram = tfidf_ngram_transformer.transform(valid_x)

# Dump the file
pickle.dump(tfidf_ngram_transformer, open("tfidf_ngram_transformer.pkl", "wb"))

```

0.4 Model 2 : Bag of word representation - Ngram level 1-3 grams

```

In [7]: # SVM Classifier on Ngram Level TF IDF Vectors
# C (penalty) = 1
# Gamma = "auto"
# Kernel : RBF (Default)

```

```

# First attempt to train the SVM classifier with max_iter = 10000 (since it performed
# but it ran endlessly for Ngram level
# Second attempt is made by reducing max_iter to 1000
results_ngram = train_model(svm.SVC(gamma = "auto", max_iter = 1000), xtrain_tfidf_ngram)
print ("SVM, N-Gram Vectors TF-IDF train accuracy: ", results_ngram[1][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF train precision: ", results_ngram[1][1])
print("")
print ("SVM, N-Gram Vectors TF-IDF train recall: ", results_ngram[1][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF train f1_score: ", results_ngram[1][3])
print("*****")
print ("SVM, N-Gram Vectors TF-IDF test accuracy: ", results_ngram[0][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF test precision: ", results_ngram[0][1])
print("")
print ("SVM, N-Gram Vectors TF-IDF test recall: ", results_ngram[0][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF test f1_score: ", results_ngram[0][3])

```

```

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning:
  % self.max_iter, ConvergenceWarning)

```

SVM, N-Gram Vectors TF-IDF train accuracy: 0.398832

SVM, N-Gram Vectors TF-IDF train precision: 0.32221193737795345

SVM, N-Gram Vectors TF-IDF train recall: 0.398832

SVM, N-Gram Vectors TF-IDF train f1_score: 0.31516400603678385

SVM, N-Gram Vectors TF-IDF test accuracy: 0.398024

SVM, N-Gram Vectors TF-IDF test precision: 0.32029641839471057

SVM, N-Gram Vectors TF-IDF test recall: 0.398024

SVM, N-Gram Vectors TF-IDF test f1_score: 0.3138662166797743

0.5 Model 3 : Bag of word representation - Ngram level 1-3 grams : Change gamma to 1

```

In [9]: # SVM Classifier on Ngram Level TF IDF Vectors
        # C (penalty) = 1
        # Gamma = 1
        # Kernel : RBF (Default)

```

```

# max_iter = 1000
results_ngram_gamma = train_model(svm.SVC(gamma = 1, max_iter = 1000), xtrain_tfidf_ng
print ("SVM, N-Gram Vectors TF-IDF train accuracy: ", results_ngram_gamma[1][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF train precision: ", results_ngram_gamma[1][1])
print("")
print ("SVM, N-Gram Vectors TF-IDF train recall: ", results_ngram_gamma[1][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF train f1_score: ", results_ngram_gamma[1][3])
print("*****")
print ("SVM, N-Gram Vectors TF-IDF test accuracy: ", results_ngram_gamma[0][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF test precision: ", results_ngram_gamma[0][1])
print("")
print ("SVM, N-Gram Vectors TF-IDF test recall: ", results_ngram_gamma[0][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF test f1_score: ", results_ngram_gamma[0][3])

```

SVM, N-Gram Vectors TF-IDF train accuracy: 0.23055466666666666

SVM, N-Gram Vectors TF-IDF train precision: 0.3166692114832456

SVM, N-Gram Vectors TF-IDF train recall: 0.23055466666666666

SVM, N-Gram Vectors TF-IDF train f1_score: 0.237200153944769

SVM, N-Gram Vectors TF-IDF test accuracy: 0.227328

SVM, N-Gram Vectors TF-IDF test precision: 0.31015521176665445

SVM, N-Gram Vectors TF-IDF test recall: 0.227328

SVM, N-Gram Vectors TF-IDF test f1_score: 0.23294178086559933

0.6 Model 4 : Bag of word representation - Ngram level 1-3 grams, (Keep gamma, other params same as model 2) : Change kernel to linear

In [8]: *# SVM Classifier on Ngram Level TF IDF Vectors*

```

# C (penalty) = 1
# Gamma = "auto"
# Kernel = Linear
# Max_iter = 1000

```

```

results_ngram_new = train_model(svm.SVC(kernel = "linear", gamma = "auto", max_iter = 1
print ("SVM, N-Gram Vectors TF-IDF train accuracy: ", results_ngram_new[1][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF train precision: ", results_ngram_new[1][1])

```

```

print("")
print ("SVM, N-Gram Vectors TF-IDF train recall: ", results_ngram_new[1][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF train f1_score: ", results_ngram_new[1][3])
print("*****")
print ("SVM, N-Gram Vectors TF-IDF test accuracy: ", results_ngram_new[0][0])
print("")
print ("SVM, N-Gram Vectors TF-IDF test precision: ", results_ngram_new[0][1])
print("")
print ("SVM, N-Gram Vectors TF-IDF test recall: ", results_ngram_new[0][2])
print("")
print ("SVM, N-Gram Vectors TF-IDF test f1_score: ", results_ngram_new[0][3])

```

SVM, N-Gram Vectors TF-IDF train accuracy: 0.37072533333333335

SVM, N-Gram Vectors TF-IDF train precision: 0.30766569065254

SVM, N-Gram Vectors TF-IDF train recall: 0.37072533333333335

SVM, N-Gram Vectors TF-IDF train f1_score: 0.3140928909583053

SVM, N-Gram Vectors TF-IDF test accuracy: 0.370408

SVM, N-Gram Vectors TF-IDF test precision: 0.3073791347288203

SVM, N-Gram Vectors TF-IDF test recall: 0.370408

SVM, N-Gram Vectors TF-IDF test f1_score: 0.313741498817925

In []: *# SAVE MODEL SO THAT IT CAN BE LOADED IN THE PREDICT SCRIPT*

```

best_svm_model = svm.SVC(gamma = "auto", max_iter = 1000)
best_svm_model.fit(xtrain_tfidf_ngram, train_y)
filename = 'best_svm_model.sav'
pickle.dump(best_svm_model, open(filename, 'wb'))

```

In []: *## FIT CALIBRATED CLASSIFIER CV TO BE ABLE TO GET PROBABILITIES*

```

filename = 'best_svm_model.sav'
loaded_model = pickle.load(open(filename, 'rb'))
svm_model = CalibratedClassifierCV(loaded_model)
svm_model.fit(xtrain_tfidf_ngram, train_y)
pickle.dump(svm_model, open("svm_model.sav", 'wb'))

```

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/model_selection/_split.py:205:

warnings.warn(CV_WARNING, FutureWarning)

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning

% self.max_iter, ConvergenceWarning)

/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning

% self.max_iter, ConvergenceWarning)


```
/Users/anjaliiverma/anaconda3/lib/python3.6/site-packages/sklearn/svm/base.py:244: ConvergenceWarning  
% self.max_iter, ConvergenceWarning)
```